

Placement with Irregularities

Mini Task Report

Florian Prott and Julian Käuser

Supervisor: Jakob Wenzel M.Sc.

Submission: 09/25/2017

Institute for Computer Engineering | Computer Systems Group | Prof. Dr.-Ing. Christian Hochberger



TECHNISCHE
UNIVERSITÄT
DARMSTADT



1 Introduction

Routing and Placement are two necessary and time-consuming steps in the FPGA synthesis toolchain. Optimizations for module design in these steps speed up or enhance the FPGA design process significantly. RapidSmith [?] allows the implementation of these steps in a Java framework. It offers an API to read in Xilinx designs and FPGA descriptions, and offers researchers the possibility to approach placement and routing problems in an object-oriented way. In RapidSoC[?], a project based on RapidSmith, a re-placement method for already placed and routed modules is integrated. Although this method works for a majority of modules, some modules routing elements cannot be re-placed due to irregular structures in the FPGA fabric. In this project, the cause for these errors is examined, and an approach to fix them is introduced.

The report is structured as following. First, the exact problem description is presented. The designed approach to the problem using isoelectric planes is described. Then, the insights gained are presented. Finally, future work is proposed, and the report is concluded.

1.1 Task Description

Since the fabric of FPGAs is in general regular, the expected error source for designs at different locations are irregularities regarding the routing elements in the fabric. In the current state, the nets of the design have to be re-routed, which is time-consuming. Since we guess that only 1-2 programmable interconnect points ("PIPs") are set falsely, a detection and repair of the conflicting nets promises faster success than a new routing. The task for this project is to determine a method which finds places where the described error occurs and is able to include the missing PIPs into the net. The proposed method makes use of the interconnect description features of RapidSmith, which operate on a mask of the FPGA fabric.

1.2 Framework and Environment

This project is embedded into the RapidSmith Java environment, version a.b.c. Classes of the RapidSoC system, especially the class `MoveModulesEverywhere`, are the basis of our work. `MoveModulesEverywhere` parses existing designs from .xdl files into the RapidSmith environment, places the included modules on any feasible destination, and generates the corresponding .xdl output files. Further processing is accomplished with Xilinx ISE [?]. Based on the output of the ISE tool xdl, the re-placed designs can be evaluated to be functional or conflicting.

All processing developed in this project is performed in a separate Java package. Every re-placed design is extracted and examined before the output step, so that the developed methods can be applied to a RapidSmith Design object.

1.2.1 Naming Conventions

Throughout this report, multiple references to nets, pins, PIPs, tiles and nodes are made. These describe the physical objects on the FPGA. Net, Pin, PIP, Tile and Node describe their object representation in RapidSmith, which is mostly equally named as the physical object.

2 General Approach to the Problem

In this chapter, our approach to find and repair conflicting spots in the netlists is described.

2.1 Routing Elements in RapidSmith

Since the RapidSmith API define multiple elements involved in the routing process, a short overview over these elements is given in the following list:

- **Wire**

A wire is the basic connection element. It is represented as an integer due to the large number of wires on an FPGA. Wires can be hard connected to other wires, which can be determined through the Device class.

- **Pin**

A Pin connects an input or output of a logic block to routing resources. It is primarily connected to one wire, and may be defined as output or input pin.

- **PIP**

A PIP object describes an active connection between two wires, which are attributes to the object. If it is contained in a net, it is switched on. Otherwise, it remains turned off.

- **Net**

A Net gathers a list of connected pins and active PIPs. It also has a defined source pin driving the net.

- **WireConnection**

A WireConnection is a wrapper class for a wire and holds the information whether this wire is only reachable if a pip is turned on. This class is used to get information about the connectable wires on the FPGA.

- **Node**

A Node is a routing object used in the router package. It can be used to describe routes of wires and PIPs in a router.

- **SinkPin**

A SinkPin refers to a switch matrix on the FPGA. Nearly all pins' wires are connected to a switch matrix, which is an interface to the other wires. SinkPins have to be viewed if a connected Pin of a wire shall be found.

- **Tile**

A Tile is one of the checkerboard-like distributed areas on an FPGA. It holds various elements from the list above and a set of logic elements.

2.2 Isoelectric Potential Search

Each design's routing in RapidSmith can be viewed as a set of Net objects. According to the RapidSmith documentation [?], these Net objects hold all used pins and pips of the net. Therefore it should be possible to reach all sink pins of a net by following wires connected to the input pin if the net is not broken. Consequently each design is broken if it is not possible to reach the output pins of an net from the input pin.

A Net in RapidSmith only contains PIP objects representing PIPs that are currently "switched on". Therefore, a net can be expanded if additional PIPs are switched on. It is theorized that it is possible to fix the broken net by switching on PIPs which reconnect the net's pins' connected wires with each other. This can only be done when the PIP will not connect the circuit to an independent third circuit.

To check whether or not a net can be traversed from its source pin to all output pins a new measurement called Potential is introduced into RapidSmith. A Potential can be seen as the isoelectric set of wires, pins and PIPs which are connected, starting from one pin. It is never possible for a Pin, PIP or Wire to be in two Potentials at the same time. It is not possible that two non-connected elements have the same Potential. This concept is named after the concept of isoelectric potentials found in classical electronics. Although two electrical potentials may be at the same numeric voltage level and still be different (not isoelectric), this is not the case for the introduced potentials.

Before any further operation on a Potential may be performed, its spatial spread must be computed. This means that, starting from a pin, each electrically connected wire, pin and pip must be found. The search for connected elements works as described in algorithm ???. It is also pointed graphically in Figure 2.1.

```
wires = {wireOf(sourcePin)};
pips = {};
pins = {sourcePin};
adjacentPIPs = {};
while new elements added do
    foreach wire i do
        foreach reacheableWire(i) do
            if !reacheable.isPIP() then
                | add reacheableWire to wires;
            end
            else if pip ∈ net.pips() then
                | add reacheableWire to wires;
                | add pip to pips;
            end
            else
                | add pip to adjacentPIPs;
            end
        end
    end
foreach wire i do
    | add connected pin to pins;
end
```

Algorithm 1: Algorithm to determine all elements on one isoelectric potential.

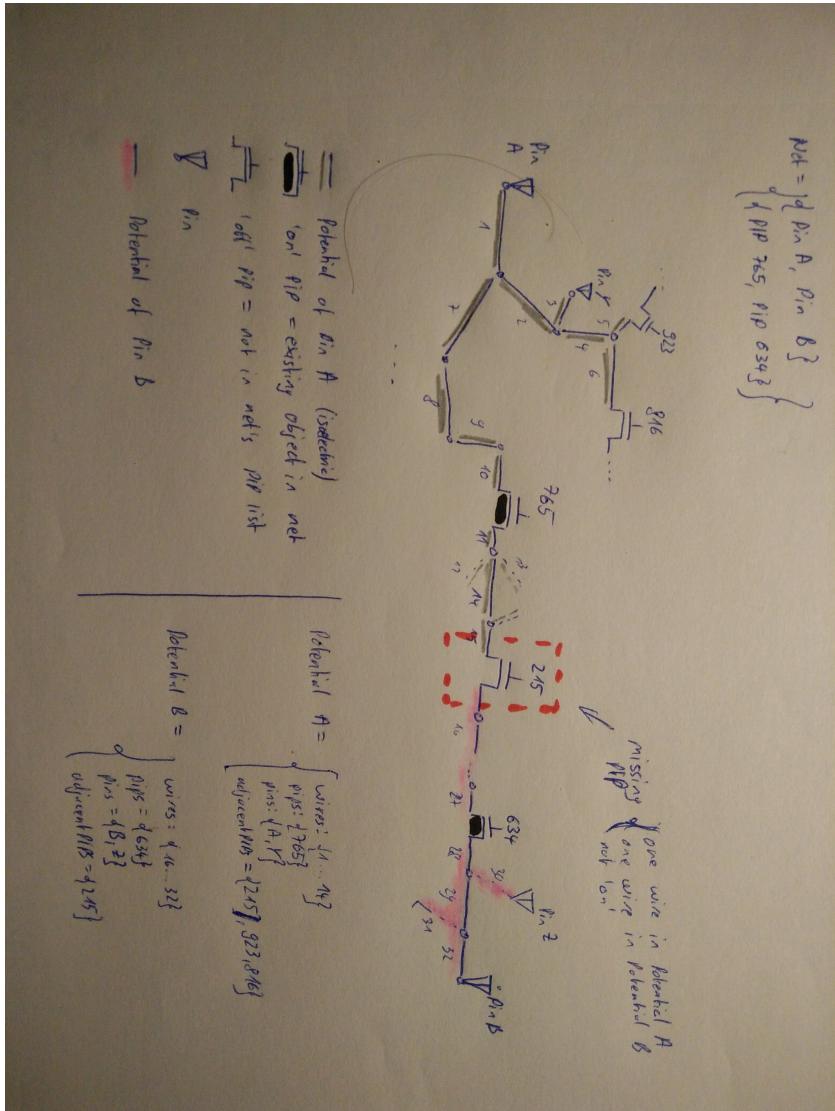


Figure 2.1: Graphical visualization of the Potential derivation.

images/brokenpotentials.pdf

Figure 2.2: A net with two Potentials which have to be connected.

Given this method, a Potential derived for one pin must always hold any other pin in the net if the net is routed correctly.

Potentials can be fused by setting PIPs. Electrically, this means that two isoelectric sets of elements are connected by a switch, and thus become one isoelectric potential. For the object representation this means that, if a pip is set, the two Potential objects have to be united.

The concept of isoelectric potentials is chosen because the search for adjacent pips wire by wire is simplified this way.

2.3 Finding and Repairing broken Nets

In order to check if a net is broken the potential of each pin is calculated and compared. If the net contains more than one unique potential then the net is broken. In that case it is necessary to reconnect the two parts (potentials) of that net. If the assumption that only one or two PIPs are missing is true, this can be done by applying the following breadth-first-search on the Potentials of pins A and B:

Let *leaves* be an ordered priority queue which holds only PIPs that do not connect to other potential besides B. *leaves* is ordered by the number of wires between the PIP and Potential A.

Put into *leaves* all adjacent PIPs of A. While *leaves* not empty do: Get first PIP of *leaves* If PIP has Potential B then done Add all adjacent PIPs to *leaves*

Figure 2.2 shows the case for one non-set PIP in a net. The net specifies two pins, A and B. For both pins, the Potential has been determined. Clearly marked are the included wires, PIPs and other pins. PIP 215 is adjacent to both potentials. The breadth-first-search finds PIP 215 to be adjacent to both pins' potentials, so this PIP has to be set (a PIP instance is created for the connected wires and added to the net). For situations where more than one PIP has to be set, the breadth-first-search steps one level down the search tree and repeats the search on every leaf. By going one level deeper, the first PIP in the priority queue has to be set.

As one can see, the breadth-first search runtime highly depends on the number of missing PIPs in the net. For every missing PIP, the number of elements to search is multiplied by the amount of adjacent PIPs to the Potential.

Although the approach described in chapter 2 seems promising, testing it for various example nets (conflicting and functional in terms of the ISE xdl output) did not lead to the expected results. Therefore, we examined the reasons, implementation and other possible error sources.

2.4 Hand Routing

The RapidSmith framework has some built-in example classes, which are designed to help the user understand the way RapidSmith works. One of these classes is the `HandRouter`, which allows the routing of a net through a console. It displays reacheable wires at the next iteration and adds PIPs as the route is chosen. We used and adapted the `HandRouter` to manually examine broken nets. Working on complete .xdl-files only, we adapted the router to work on single nets, and inserted it after the potential derivation.

Using the `HandRouter` reveals several error sources, but also that each connection reacheable from each pin is correctly integrated in the pin's `Potential` instance.

3 Insights

In the process of the task, we found out that... X Y Z doesn't work as expected

3.1 Possible Solutions

For the encountered problems, we propose the following approaches

4 Conclusion

In conclusion it seems that there is currently no possibility to determine the exact properties of a net with RapidSmith based on the "xdl" files. Therefore there is no way to determine which net might be broken or fix the broken design. During our experiments with the RapidSmith hand router we started to suspect that our database might be too small. Most nets, even those of not broken designs e.g. designs that can be used without a problem by the ISE tool-chain, contain no wires or other elements connecting the pins with each other. The few Nets that are connected are usually of a rather simplistic manner. Alternatively there is the slim possibility of bugs within RapidSmith. While we did not encounter any evidence for bugs we must note that the RapidSmith version used in this Project showed some differences to the publicly available version on <http://rapidsmith.sourceforge.net/>. While most of these differences seem to be improvements, for example the wireEnumerator was upgraded by one version, it should also be noted that the GUI seems to be missing files.

All things considered we recommend to use the algorithmic approach of this work with a bigger dataset or an alternative tool.

<http://rapidsmith.sourceforge.net/papers/Nelson-FPL11-Presentation.pdf>

We suspect As shown in REF are X to Y big and ours are simply

the task was to - examine re-placed designs' net lists

- find out if something does not work
- expected: only one or two PIPs are missing due to irregular structure of FPGA
- if problem is encountered, find rule/mask based method to fix these missing pips

what we did was:

- handle net by net
- for each Pin, determine all elements of the isoelectric it is connected to (class potential)
- checking method: if $\text{potential}(\text{source}) \neq \text{potential}(\text{sink}) \rightarrow \text{broken}$
- for broken nets:
 - search isoelectric of source, sink for adjacent, non-set PIPs (intersection of sets). if $\neq \text{empty set}$, activate this pip
- why it did not work:
 - even for correct (in terms of make file) nets, the method fails

- hand routing (function of RapidSmith framework) neither does

- we suspect missing information in fabric;

- according to the hand router, the method of "isoelectric search"

results in correct sub nets; might be useful for later work