

Placement with Irregularities

Mini Task Report

Florian Prott and Julian Käuser

Supervisor: Jakob Wenzel M.Sc.

Submission: 09/25/2017

Institute for Computer Engineering | Computer Systems Group | Prof. Dr.-Ing. Christian Hochberger



TECHNISCHE
UNIVERSITÄT
DARMSTADT



1 Introduction

Routing and Placement are two necessary and time-consuming steps in the FPGA synthesis toolchain. These processes can sometimes be optimized or speed up to enhance the FPGA design process significantly. The RapidSmith[3] Java framework is a freely available tool that helps with these steps. It offers an API to read in Xilinx designs and FPGA descriptions, and offers researchers the possibility to approach placement and routing problems in an object-oriented way. In RapidSoC[4], a project based on RapidSmith, a re-placement method for already placed and routed modules is integrated. Although this method works for the majority of modules, some modules routing elements cannot be re-placed due to irregular structures in the FPGA fabric. In this project, the cause for these errors is examined, and an approach to fix them is introduced.

The report is structured as following: First, the exact problem description is presented. The designed approach to the problem using isoelectric wires is described. After that, the insights gained are presented. Finally, future work is proposed, and the report is concluded.

1.1 Task Description

To speed up the FPGA development cycle it can be preferable to move a fixed design on an FPGA from one location to another without the process of rerouting the design. If the FPGA is regular, like the first iterations of FPGAs generally were, this process can be done without any problem. Unfortunately, FPGA vendors introduced some irregularities in their designs to improve the efficiency of the FPGA and lower the costs. Due to this it is possible that an FPGA design that was moved on the FPGA is no longer operational after the re-placement. The task of this project is to check if a design is broken and repair the design if possible. It is suspected that a corrupted part of the FPGA only contains a small amount of irregularities (about 1-2 programmable interconnect points ("PIPs")) and detection and repair of the conflicting nets is faster than new routing. This work makes use of the interconnect description features of RapidSmith, which operate on a mask of the FPGA fabric.

1.2 Framework and Environment

This project is embedded into the RapidSmith Java environment, provided by the RS lab. It uses classes of the RapidSoC system. Especially the class `MoveModulesEverywhere` acts as the basis of our work. `MoveModulesEverywhere` parses existing designs from .xdl files into the RapidSmith environment, places the included modules on any feasible destination, and generates the corresponding .xdl output files. Further processing is accomplished with Xilinx ISE [1]. Based on the output of the ISE tool xdl, the re-placed designs can be evaluated to be functional or conflicting.

All processing developed in this project is performed in a separate Java package. Every re-placed design is extracted and examined before the output step, so that the developed methods can be applied to a RapidSmith Design object.

1.2.1 Naming Conventions

Throughout this report, multiple references to nets, pins, PIPs, tiles and nodes are made. These describe the physical objects on the FPGA. `Net`, `Pin`, `PIP`, `Tile` and `Node` describe their object representation in RapidSmith, which is mostly equally named as the physical object.

2 General Approach to the Problem

In this chapter, our approach to find and repair conflicting spots in the netlists is described.

2.1 Routing Elements in RapidSmith

Since the RapidSmith API define multiple elements involved in the routing process, a short overview over these elements is given in the following list:

- **Wire**
A Wire is the basic connection element. It is represented as an integer due to the large number of wires on an FPGA. Wires can be hard connected to other wires, which can be determined through the Device class.
- **Pin**
A Pin connects an input or output of a logic block to routing resources. It is primarily connected to one wire, and may be defined as output or input pin.
- **PIP**
A PIP object describes an active connection between two wires, which are attributes to the object. If it is contained in a net, is is switched on. Otherwise, it remains turned off.
- **Net**
A Net gathers a list of connected pins and active PIPs. It also has a defined source pin driving the net.
- **WireConnection**
A WireConnection is a wrapper class for a wire and holds the information whether this wire is only reachable if a PIP is turned on. This class is used to get information about the connectable wires on the FPGA.
- **Node**
A Node is a routing object used in the router package. It can be used to describe routes of wires and PIPs in a router.
- **SinkPin**
A SinkPin refers to a switch matrix on the FPGA. Nearly all pins' wires are connected to a switch matrix, which is an interface to the other wires. SinkPins have to be viewed if a connected pin of a wire shall be found.
- **Tile**
A Tile is one of the checkerboard-like distributed areas on an FPGA. It holds various elements from the list above and a set of logic elements.

2.2 Isoelectric Potential Search

Each design's routing in RapidSmith can be viewed as a set of Net objects. According to the RapidSmith documentation [2], these Net objects hold all used pins and pips of the net. Therefore it should be possible to reach all sink pins of a net by following wires connected to the input pin if the net is not broken. Consequently each design is broken if it is not possible to reach the output pins of a net from the input pin.

A Net in RapidSmith only contains PIP objects representing PIPs that are currently "switched on", but no Wires. Therefore, a net can be expanded if additional PIPs are switched on. It is theorized that it is possible to fix the broken net by switching on PIPs which reconnect the net's pins' connected wires with each other. This can only be done when the PIP will not connect the circuit to an independent third circuit.

To check whether or not a net can be traversed from its source pin to all output pins a new measurement called Potential is introduced into RapidSmith. A Potential can be seen as the isoelectric set of wires, pins and PIPs which are connected, starting from one pin. It is never possible for a Pin, PIP or Wire to be in two Potentials at the same time. It is not possible that two non-connected elements have the same Potential. This concept is named after the concept of isoelectric potentials found in classical electronics. Although two electrical potentials may be at the same numeric voltage level and still be different (not isoelectric), this is not the case for the introduced potentials.

Before any further operation on a Potential may be performed, its spatial spread must be computed. This means that, starting from a pin, each electrically connected wire, pin and pip must be found. The search for connected elements works as described in algorithm 1. It is also pointed graphically in Figure 2.1. Given this method, a Potential derived for one pin must always hold any other pin in the net if the net is routed correctly.

Potentials can be fused by setting PIPs. Electrically, this means that two isoelectric sets of elements are connected by a switch, and thus become one isoelectric potential. For the object representation this means that, if a PIP is set, the two Potential objects have to be united.

In order to keep track of the switchable connections to other potentials, each potential has a set of adjacent PIPs. When the potential is created and its spread is computed, every PIP which is connected to any wire of the potential, but not specified to be switched on by the net, is added to the adjacent PIP set. Hence, all PIPs which can be reached from the pins in the Potential are stored.

The concept of isoelectric potentials is chosen because the search for adjacent pips wire by wire is simplified this way.




Figure 2.1: Graphical visualization of the Potential derivation. The wire segments and elements are enumerated in the sequence of integration. PIP IDs are in parentheses.

```

wires = {wireOf(sourcePin)};
pips = {};
pins = {sourcePin};
adjacentPIPs = {};
while new elements added do
    foreach wire i do
        foreach reacheableWire(i) do
            if !reacheable.isPIP() then
                | add reacheableWire to wires;
            end
            else if pip ∈ net.pips() then
                | add reacheableWire to wires;
                | add pip to pips;
            end
            else
                | add pip to adjacentPIPs;
            end
        end
    end
end
foreach wire i do
    | add connected pin to pins;
end

```

Algorithm 1: Algorithm to determine all elements on one isoelectric potential.

2.3 Finding and Repairing broken Nets

In order to check if a net is broken the potential of each pin is calculated and compared. If the net contains more than one unique potential then the net is broken. In that case it is necessary to reconnect the two parts (potentials) of that net. If the assumption that only one or two PIPs are missing is true, this can be done by applying the following breadth-first-search on the Potentials of pins A and B:

There is a priority queue *leaves* which holds only PIP-elements that do not connect to other Potentials besides B. Each of these elements has a parent PIP which introduced the PIP into the queue. *Leaves* is ordered by the number of parents between the PIP and Potential A in ascending order. Algorithm 2 points out the search on that priority queue.

Figure 2.2 shows the case for one non-set PIP in a net. The net specifies two pins, A and B. For both pins, the Potential has been determined. Clearly marked are the included wires, PIPs and other pins. PIP 215 is adjacent to both potentials. The breadth-first-search finds PIP 215 to be adjacent to both pins' potentials, so this PIP has to be set (a PIP instance is created for the connected wires and added to the net). For situations where more than one PIP has to be set, the breadth-first-search steps one level down the search tree and repeats the search on every leaf. By going one level deeper, the first PIP in the priority queue has to be set.

As one can see, the breadth-first search runtime highly depends on the number of missing PIPs in the net. For every missing PIP, the number of elements to search is multiplied by the amount of adjacent PIPs to the Potential.

```

add all adjacent PIPs of A to leaves;
while leaves.size() > 0 do
  PIP pip = leaves.pop();
  if Potential B contains pip then
    | return pip; // we can determine the connection trough the parent PIPs;
  end
  foreach PIP p ∈ adjacentPIPs(potential of pip) do
    | if p ≠ pip.parent then
    | | leaves.push(p); //if p is of an other potential than B we cannot add it ;
    | end
  end
end
end

```

Algorithm 2: Exemplary algorithm to determine missing PIPs by breadth-first-search

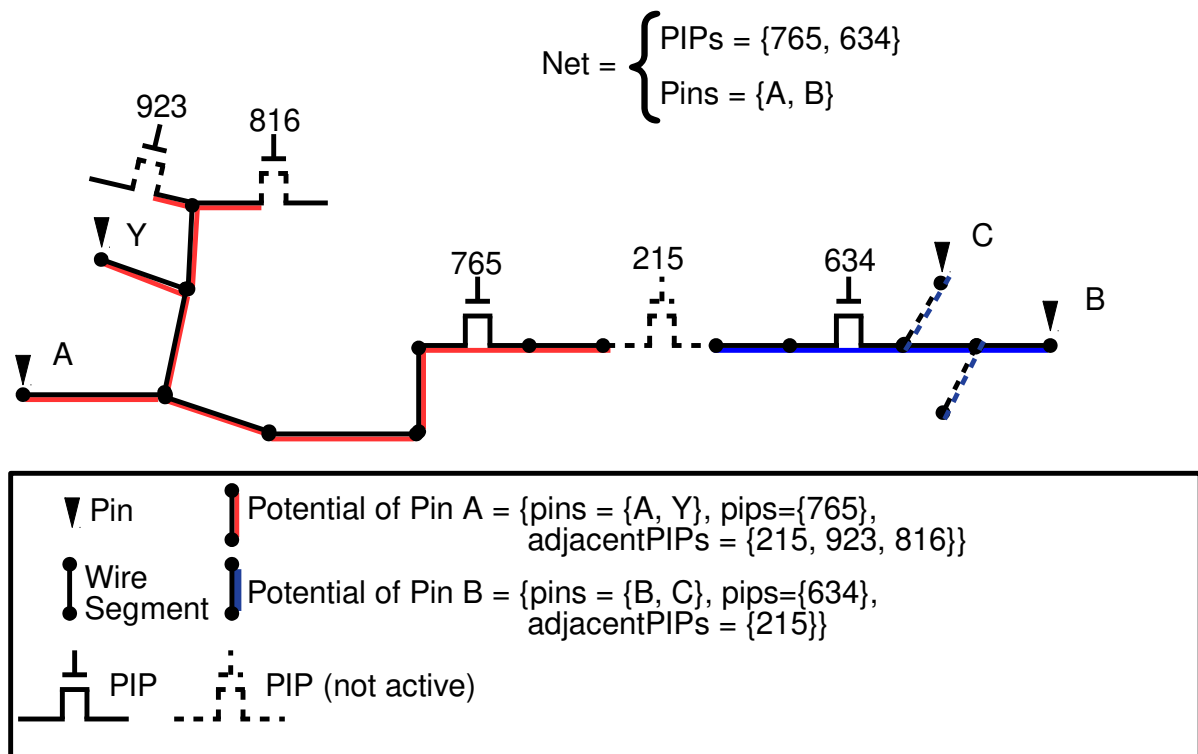


Figure 2.2: A net with two Potentials which have to be connected.

3 Insights

Although the approach described in chapter 2 seems promising, testing it with different designs and for various example nets (both conflicting and functional in terms of the ISE xdl output) did not lead to the expected results. Instead of finding exactly one potential per net, the most nets contain more than one potential. This behavior occurs regardless whether or not the design can be used by ISE and therefore has not to be rerouted. Therefore, we examined the reasons, implementation and other possible error sources that may lead to this behavior.

3.1 Hand Routing

The RapidSmith framework has some built-in example classes, which are designed to help the user understand the way RapidSmith works. One of these classes is the `HandRouter`, which allows the routing of a net through a console. It displays reachable wires at the next iteration and adds PIPs as the route is chosen. We used and adapted the `HandRouter` to manually examine broken nets. Working on complete .xdl-files only, we adapted the router to work on single nets, and inserted it after the potential derivation.

Using the `HandRouter` reveals several error sources, but also that each connection reachable from each pin is correctly integrated in the pin's `Potential` instance. Therefore, we could prove that, based on the given nets, the concept of isoelectric potentials is implemented correctly. Nevertheless, the potentials and the routing of working designs were still not as expected.

Further analysis of the reachable routing elements from every pin revealed that the assumption of only 1 or two missing PIPs between the isoelectric potentials of the sink and source pin evaluated to be wrong. In the most examined cases, at least four or more PIPs not specified by the net have to be additionally switched on. While some of the wire segments where additional PIPs are necessary only feature one or two switchable PIPs, others had a much longer list.

Therefore we were not able to determine which net is broken and therefore were not able to apply the breadth-first-search algorithm. In order to fix a design we would have to apply the algorithm on nearly every net of the design. The breadth-first-search is suitable for search trees of one or two levels, in order to fulfill the purpose of this work, to save time during development, it is advisable to use a router able to solve the problem of many possible connections and a larger number of PIPs.

3.2 Possible Reasons for Unexpected Behaviour

The number of missing PIPs leads us to the assumption that the errors in re-placed designs or nets are not completely caused by small irregularities in the FPGA fabric. RapidSmith features a coordinate-based location description of nearly all elements of the FPGA. While this is useful for the majority of applications, some details may be hidden in this abstraction. Not every cartesian coordinate perfectly translates into the actual position of the element on the FPGA, so that regular structures in terms of the cartesian coordinates do not necessarily imply regular structures in the fabric. Hence, the assumption of conflicting re-placed nets caused by slight irregularities could not be proven.

In addition to the larger-than-expected irregularities, there are other error sources in RapidSmith. We cannot determine the exact source, but we have reason to assume that certain information on the FPGA are either missing or not correct. This assumption is made due to the fact even for functional nets (in terms of the xdl conversion), some of the corresponding RapidSmith Nets were not routed correctly (i.e. the source was reachable from all sink pins).

4 Conclusion

The goal of this project was to determine and repair nets which are corrupted after a module re-placement due to slight irregularities in the FPGA fabric. This goal could partially be met. The analysis of corrupted nets is functional. Although there is a method to fix the errors, we suggest a re-routing of the nets, because our solution is not applicable for the encountered error sources.

Our approach is based on isoelectric potentials, on which a breadth-first search finds missing interconnect elements. As pointed out in chapters 2 and 3, the basic approach has correct results. Nevertheless, the errors encountered in the conflicting nets have shown to be of a kind which is not applicable to the breadth-first-search as solution.

Reviewing the results proves that there is either a problem with RapidSmith or the assumption of only few missing PIPs is wrong. In case that RapidSmith can be ruled out as an error source, we suggest the use of a dedicated router to repair the corrupted nets, because it is suited well for connecting elements with a higher amount of interconnect points between them.

It should be noted that there is the slim possibility of bugs within RapidSmith. While we did not encounter any evidence for bugs, our experiments with the hand router showed ambiguous results regarding the functional nets. We must note that the RapidSmith version used in this project showed some differences to the publicly available version on <http://rapidsmith.sourceforge.net/>. While most of these differences seem to be improvements, for example the wireEnumerator was upgraded by one version, it should also be noted that the GUI seems to be missing files.

All things considered, we recommend to use the algorithmic approach of this work with as a starting point for further examinations of the topic. For the meanwhile, we recommend to use a re-routing in case of error, because it is probably better than the suggested approach. For future considerations, other approaches may provide better results.

All things considered, we recommend to use the algorithmic approach of this work with a bigger dataset or an alternative tool. For the meanwhile, we recommend to use a re-routing in case of error, because it is probably better than the suggested approach. For future considerations, other approaches may provide better results.

5 Bibliography

- [1] Xilinx ise integrated synthesis environment, 10 2013.
- [2] Jaren Lamprecht Philip Lundrigan Brent Nelson Brad Hutchings Christopher Lavin, Marc Padilla and Michael Wirthlin. A library for low level manipulation of partially placed and routed fpga designs. Technical report, NSF Center for High Performance Reconfigurable Computing (CHREC), Department of Electrical and Computer Engineering, Brigham Young University, <http://rapidsmith.sourceforge.net/doc/TechReportAndDocumentation.pdf?format=raw>, 1 2014.
- [3] Christopher Lavin, Marc Padilla, Philip Lundrigan, Brent Nelson, and Brad Hutchings. Rapid prototyping tools for fpga designs: Rapidsmith. In *Field-Programmable Technology (FPT), 2010 International Conference on*, pages 353–356. IEEE, 2010.
- [4] Jakob Wenzel and Christian Hochberger. Rapidsoc: short turnaround creation of fpga based socs. In *Rapid System Prototyping (RSP), 2016 International Symposium on*, pages 1–7. IEEE, 2016.